

Software for Mind Control

Ben Pearre

March 15, 2016
(Rev. -2)

Outline

Introduction

- Goals

- Previous work

Materials

- Electrodes

- Stimulator

Experiments

- Chronic recording

- Stimulation

- Current Steering

The Future!

- Ideas

- Electrodes


- Gradient estimation

Controlling behaviour with chronically implanted electrodes

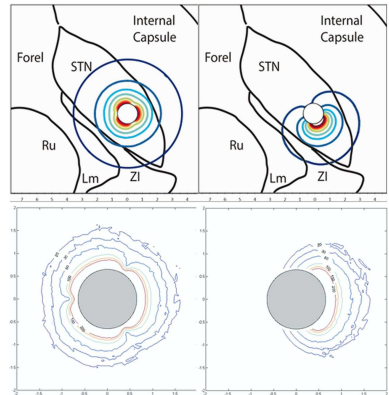
- ▶ Long-term stable recording
- ▶ Safe stimulation (despite small surface area)
- ▶ Goal-directed modification of behaviour
 - ▶ Optimise stimulation parameters
 - ▶ Learn to produce desired output
 - ▶ Safety constraints. . .

Pollo... Schüpbach [2014]

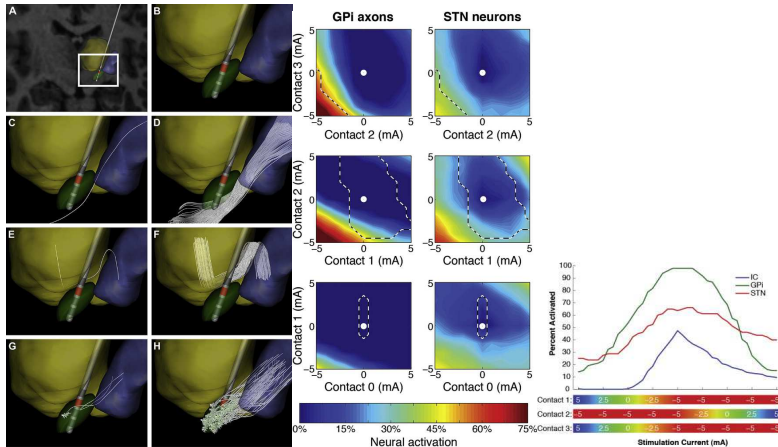
Slightly directional electrode



Finite element analysis



Chaturvedia, Foutza, McIntyre [2015]



Previous work

Current Steering using finite element analysis

- ▶ Fast
- ▶ Somewhat effective
- ▶ Coarse-grained
- ▶ All models are wrong. Some models are useful.


Recording for “closed-loop” therapy



- ▶ Good luck


Electrodes

Chronic high-count
high-impedance. . .

- ▶ Carbon fibres
- ▶ Silicon carbide
- ▶ Optical. . . ?



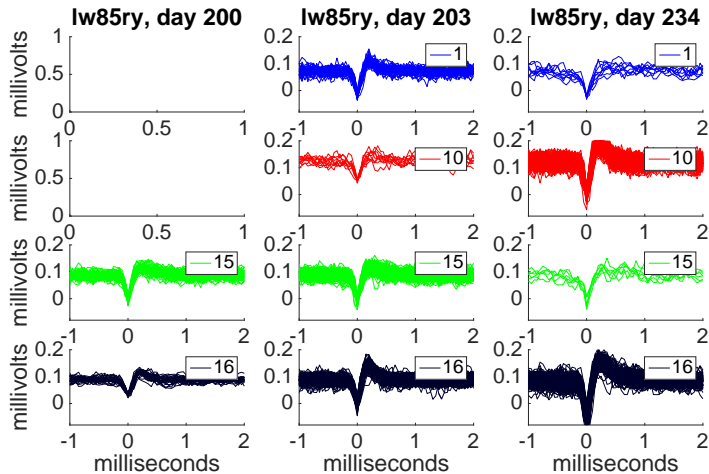





Plexon stimulator

- ▶ 16 channels
- ▶ Current-controlled
- ▶ Externally triggered
- ▶ Arbitrary pulse waveforms
- ▶ Resolution: $30 \text{ nA} \times 1\mu\text{s}$
- ▶ Matlab API
- ▶ Reprogramming time $\approx 0.2\text{s/channel}$
- ▶ Voltage monitoring is expensive!

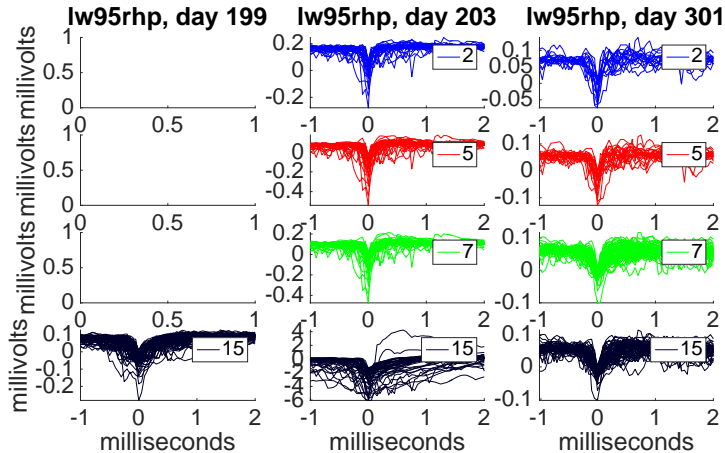
Recording — bare carbon in X




Questionable recording — IrO_2 in X



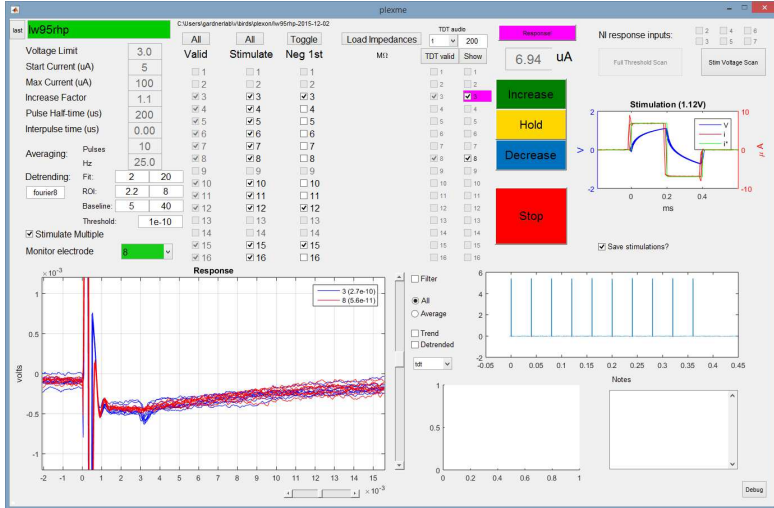
Questionable recording — IrO_2 in X





Antedromic HVC \leftarrow X response




Antedromic HVC ← X response




Combinatorial optimisation



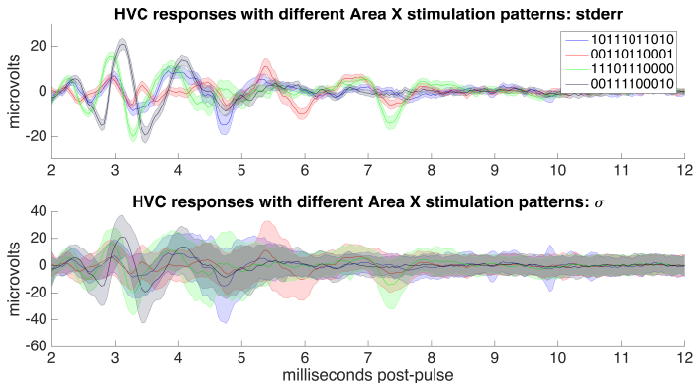
Voltage minimisation



Voltage minimisation



Response shaping in HVC



Policy optimisation

Criteria


- ▶ See response
- ▶ Maximise response
- ▶ Minimise voltage
- ▶ Separate responses
- ▶ Directed change to song
 - ▶ Acute
 - ▶ Chronic

Policy outputs

- ▶ Pulse train timing
- ▶ Channel timing
- ▶ Arbitrary pulse shape
- ▶ *Optical!*

Policy inputs

- ▶ Vocalisation
- ▶ Neural activity
- ▶ Other motor output?



Gradient estimation: eR / Stochastic Approximation

Policy:

$$\pi(s, u; \theta) = \Pr(u|s; \theta)$$

Gradient:

$$\widehat{\nabla_{\theta} J(\theta)} = \hat{g}_{\theta} = \left\langle \left(\sum_{k=0}^H \nabla_{\theta} \log \Pr(u_k | s_k; \theta) \right) \cdot (r - b) \right\rangle$$

Learning:

$$\theta_{e+1} = \theta_e + \alpha \frac{\nabla_{\theta} J(\theta)}{|\nabla_{\theta} J(\theta)|}$$

Gradient estimation: eR / Stochastic Approximation

Reward:

$$r(d, m) = - \left(d + \eta \sum_{j=1}^n \left(\max \left[0, \left(\frac{\mu D}{m_j} \right)^2 - 1 \right] \right) \right)$$

Eligibility:

$$\begin{aligned} u &= \theta + \mathcal{N}(0, \Sigma) \\ \nabla_{\theta} \log \Pr(u|s; \theta) &= \frac{1}{2} \left(\Sigma^{-1} + \Sigma^{-1 T} \right) (u - \theta) \end{aligned}$$