

20 μm
└──┘

Mag = 197 X

WD = 7.8 mm

EHT = 3.00 kV

Aperture Size = 30.00 μm


Signal A = SE2

Stage at T = 37.7 °

Signal B = SE2

Date :29 Mar 2017





10 μm

Mag = 1.79 K X

EHT = 3.00 kV

Signal A = SE2

Signal B = SE2


WD = 7.8 mm

Aperture Size = 30.00 μm

Stage at T = 37.7 °

Date :29 Mar 2017





30 μ m
|————|

Mag = 223 X

WD = 8.0 mm

EHT = 3.00 kV

Aperture Size = 30.00 μ m

Signal A = SE2

Stage at T = 46.7 °

Signal B = SE2

Date :29 Mar 2017



周文敘重屏會棋圖真蹟

敘甚善。有紹子容。四月間。有義興
任事。以遂從。金焦。渡。保。難。隨。之。能。保。可
關。公。亦。欲。歸。之。有。此。中。好。筆。可。各。州。之
我。各。保。其。封。同。其。大。道。也。因。此。事。
乾。隆。十。年。丁。卯。月。上。嘉。吉。辰。御。筆。











Google DeepMind Challenge Match

R - 15 March 2017



AlphaGo




Lee Sedol



Go tutorial...

► The Way To Go

Machine play with Minimax:



How big is the game tree?

Branching factor: At each position, n possible moves

Length: game lasts m moves

n^m possible games

Size of games

| Game | Branching factor | Length | $\log_{10} (\# \text{ Games})$ |
|---------------------------|------------------|--------|--------------------------------|
| Tic-Tac-Toe | 4 | 9 | 5 |
| Connect Four | 4 | 36 | 21 |
| Checkers (8×8) | 2.8 | 70 | 31 |
| Chess | 35 | 70 | 108 |
| Backgammon | 250 | 55 | 132 |
| Carcassonne | 55 | 71 | 195 |
| Go | 250 | 150 | 360 |

Reducing the computational complexity

Length:

- ▶ Estimate the value of non-terminal states

Branching factor:

- ▶ α - β Pruning
 - ▶ $O(\sqrt{n^m})$ if moves are sorted
- ▶ Monte Carlo Tree Search of *promising moves*

Estimating the value of non-terminal states

Checkers:

- ▶ Material
- ▶ Position of non-kings

Chess:

- ▶ Material
- ▶ Mobility
- ▶ Control centre of board


Go:

- ▶ “Strength” / “Influence”
 - ▶ Um...


Intuition!

Monte Carlo Tree Search


Selection




Expansion



Simulation



Backpropagation



Selecting a move for roll-out

Early:

- High prior probability?
- Low visit count

Later:

- Success rate

Return

The most visited move


Ignoring questionable moves

Only examine promising moves.

Intuition

In the beginner's mind there are many possibilities. In the expert's mind there are few. –Shunryu Suzuki


Neural Networks



- ▶ Input pattern ξ
- ▶ Weights W
- ▶ $y = W_1 \tanh(W_0 \xi)$
- ▶ W_i are trained with backpropagation of errors

Convolutional Neural Networks

Inspired by Hubel & Wiesel 1968?



► LeNet

Supervised learning of policy networks: $p_{\sigma}(a|s)$

Predict the expert's move:

- ▶ Given board state s , predict probability of move a
- ▶ 13 convolutional layers
 - ▶ Input: 48 features
 - ▶ Layer 1: 192 of 5×5
 - ▶ Layers 2–12: 192 of 3×3
 - ▶ Rectifier nonlinearities
- ▶ Training: 30 million positions from KGS
- ▶ 57% accurate
- ▶ Evaluation time: 3 ms

$$\Delta\sigma \propto \frac{\partial \log p_{\sigma}(a_t|s_t)}{\partial \sigma}$$

Supervised learning of policy networks II: $p_{\pi}(a|s)$

Predict the expert's move (faster):

- ▶ Given board state s , predict probability of move a
- ▶ Small pattern features
- ▶ Training: 8 million positions from Tygem (?)
- ▶ 24% accurate
- ▶ Evaluation time: $2 \mu s$

Why? To be continued...

Reinforcement learning of policy networks: $p_\rho(a|s)$

Don't predict expert moves: Win!

- ▶ Start with the supervised move predictor: $\rho \leftarrow \sigma$
- ▶ Training: self-play vs. a previous iteration of ρ
- ▶ Result of game is z

$$\Delta\rho \propto \frac{\partial \log p_\rho(a_t|s_t)}{\partial \rho} z$$

p_ρ won 80% of its games vs. p_σ

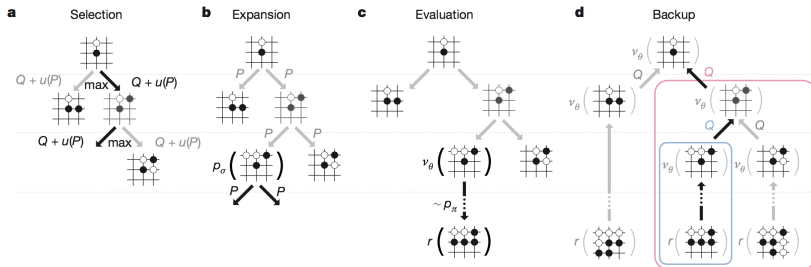
Reinforcement learning of value networks: $v^p(s)$

Predict the likelihood of winning from any board position

- ▶ $v^p(s) = \mathbb{E}[z | s = s_t, a_{t...T} \sim p]$
- ▶ Similar architecture to p_ρ , but outputs scalar v
- ▶ Trained from 30 million samples from self-play by $p_\rho(s, a)$
 - ▶ Decorrelate sequences of moves

$$\Delta\theta \propto \frac{\partial v_\theta(s)}{\partial \theta} (z - v_\theta(s))$$

Monte Carlo Tree Search II



Selecting a move for roll-out

Early:

- ▶ High prior probability from $p_\sigma(s, a)$!
- ▶ Low visit count

Later:




- ▶ Success rate

Return

The most visited move

Choosing a move




- ▶ Evaluate potential new position s' in 2 ways:
 - ▶ $v_{\theta}(s')$
 - ▶ Monte Carlo rollout (using the fast policy p_{π})
 - ▶ Initialise MCTS priors $\forall a$ from $p_{\sigma}(s', a)$ + exploration term
- ▶ Combine the evaluations (λ).

a Value network**b** Tree evaluation from value net**c** Tree evaluation from rollouts

Evaluation of all successors s of the root position s , using the value network $v_\theta(s)$; estimated winning percentages are shown for the top evaluations.

Action values $Q(s, a)$ for each edge (s, a) in the tree from root position s ; averaged over value network evaluations only ($\lambda = 0$).

Action values $Q(s, a)$, averaged over rollout evaluations only ($\lambda = 1$).

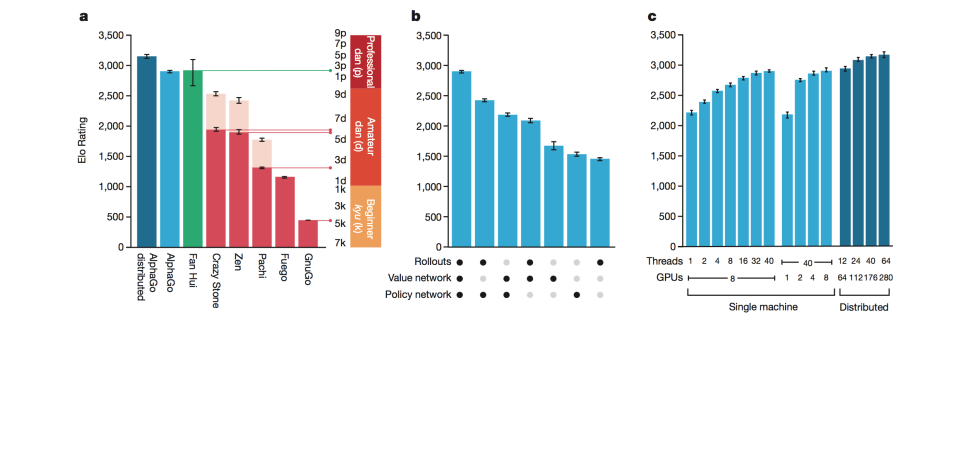
d Policy network**e** Percentage of simulations**f** Principal variation

Move probabilities
directly from the SL
policy network,
 $p_{\sigma}(a|s)$; reported as a
percentage (if above
0.1%)

Percentage frequency
with which actions
were selected from
the root during
simulations.

AlphaGo's move, and
the most likely
continuation of the
game.

Results



Victory!

- ▶ Beat Fan Hui 2p 5-0
- ▶ Beat Lee Sedol 9p 4-1
 - ▶ One of 'em!
- ▶ 10^3 times fewer position evaluations than Deep Blue!
 - ▶ Learns to aggressively examine only a few promising moves
 - ▶ Learns intuition for values of intermediate positions
 - ▶ Plays like a human?
- ▶ “AlphaGo’s play makes us feel free, that no move is impossible. Now everyone is trying to play in a style that hasn’t been tried before.” –Zhou Ruiyang 9p

Future work

- ▶ Ke Jie 9p
 - ▶ May 23–27

